

Introduction to Python Programming


(for Reinforcement Learning)

Table of Contents

1. Personal Introduction
2. Installing Python
3. Introduction to Python
 - a. Why Python for Reinforcement Learning (RL)?
 - b. Control Flow
 - c. Functions
 - d. Object Oriented Programming (OOP)
4. Python for (Deep) Reinforcement Learning
 - a. Numpy
 - b. TensorFlow / Pytorch
 - c. OpenAI Gymnasium
5. How to Debug on the Fly

Personal Introduction

- Senior at UCF studying Computer Engineering
- Co-authored Deep RL research with Professor Enyioha
- Done Deep Learning and robotics research and engineering at MIT and NASA for internships



Installing Python

Setup Instructions (Windows) - Python Install

1. Go to the official download page:
 - a. <https://www.python.org/downloads/release/python-31011/>
 - b. Scroll down to Files, and download:
 - i. Windows installer (64-bit) — Windows installer (64-bit) (.exe)
2. Run the installer:
 - a. Go to File Explorer and open "Downloads"
 - b. Double-click "python-3.10.11-amd64.exe"
 - c. Uncheck "Use Admin Privileges"
 - d. Check "Add Python 3.10 to PATH"
 - e. Click "Install Now"
 - f. If this does not work or an error comes up, let me know

Setup Instructions (Windows) - Environment Setup

1. Verify Python

- a. Open Command Prompt
- b. Run: `python --version`
 - i. Should say "Python 3.10.11"

2. Make requirements.txt

- a. Open Notepad
- b. Write these two lines EXACTLY:
`gymnasium[classic-control,atari]`
`stable-baselines3[extra]`
- c. Save the file, call it "requirements.txt" and save it to your Downloads folder

3. Install Libraries

- a. `python -m pip install -U pip`
- b. `cd Downloads`
- c. `python -m pip install -r .\requirements.txt`

Setup Instructions (Linux)

1. Install Python 3.10.11

- a. `sudo apt update && sudo apt upgrade -y`
- b. `sudo apt install software-properties-common -y`
- c. `sudo add-apt-repository ppa:deadsnakes/ppa`
- d. `sudo apt install python3.10`


2. Setup Environment

- a. `python3.10 -m venv cca-env`
- b. `source cca-env/bin/activate`
- c. `pip install --upgrade pip`
- d. `pip install -r requirements.txt`

Introduction to Python Programming

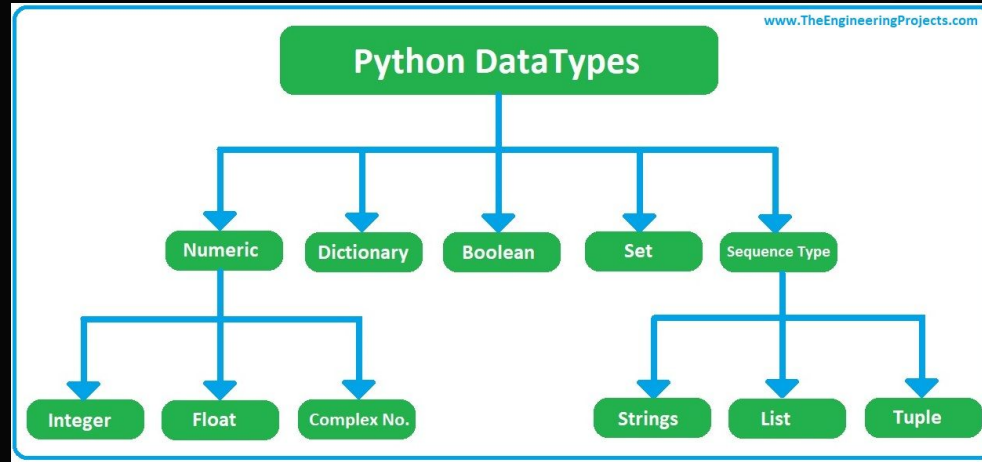
Why Python for Reinforcement Learning?

- Python is one of the easiest programming languages to learn
- Decades of development in machine learning in Python means there are plenty of tools readily available to use in Python
- Researchers and Engineers don't have to code everything from scratch if we use these tools



Relevant XKCD Comic

Variables



- All variables can be any kind of data type in Python

Operators

How we change the value of variables

- Arithmetic
 - + (Addition)
 - - (Subtraction)
 - * (Multiplication)
 - / (Division)
- Logical
 - "and"
 - "or"
 - "Not"
- Comparison
 - $a == b$ (are a and b equal?)
 - $a != b$ (are a and b not equal?)

Control Flow - If Statements

- If Statements
 - “If a equals 0, print “0”.
Else if a equals 2, print “2”.
Else print nothing”

```
a = 1
if (a == 0):
    print(0)
elif (a == 2):
    print(2)
else:
    print()
```

Control Flow - Loops

- For Loops
 - “For each element in this sequence, do this”

```
number_sequence = [1, 2, 3, 4, 5]
for number in number_sequence:
    print(number, end=", ")
# Output:
# 1, 2, 3, 4, 5,
```

- While Loops
 - “Loop until this condition is not true”

```
number = 0
while (number != 5):
    number += 1
    print(number, end=", ")
print("\nexiting, number now equals 5")
# Output:
# 1, 2, 3, 4, 5,
# exiting, number now equals 5
```

Functions

- This is how programmers abstract away large amounts of code
 - If you want to run the same 10 lines of code 3 times, you don't want to write the 10 lines 3 times. A function lets you write it once and call it 3 times.
- Exact same as the mathematical function notation
 - $y = f(x)$, where f is the function
- Syntax:
 - **'def'**: declare function definition
 - **'bubble_sort'**: function name
 - **'arr'**: the only argument for this function, there can be more
 - **'return'**: what value to return ('y' in $y = f(x)$)

```
def bubble_sort(arr):  
    n = len(arr)  
    for i in range(n):  
        # Last i elements are already sorted  
        for j in range(0, n - i - 1):  
            # Swap if the current element is greater than the next  
            if arr[j] > arr[j + 1]:  
                arr[j], arr[j + 1] = arr[j + 1], arr[j]  
    return arr
```

Functions (Example)

Using Functions:

```
def bubble_sort(arr):
    n = len(arr)
    for i in range(n):
        # Last i elements are already sorted
        for j in range(0, n - i - 1):
            # Swap if the current element is greater than the next
            if arr[j] > arr[j + 1]:
                arr[j], arr[j + 1] = arr[j + 1], arr[j]
    return arr

num1 = [5, 1, 4, 2, 8]
num2 = [6, 3, 10, 1, 2]
num3 = [2, 3, 1, 8, 4]

print(bubble_sort(num1)) # Output: [1, 2, 4, 5, 8]
print(bubble_sort(num2)) # Output: [1, 2, 3, 6, 10]
print(bubble_sort(num3)) # Output: [1, 2, 3, 4, 8]
```

Not Using Functions:

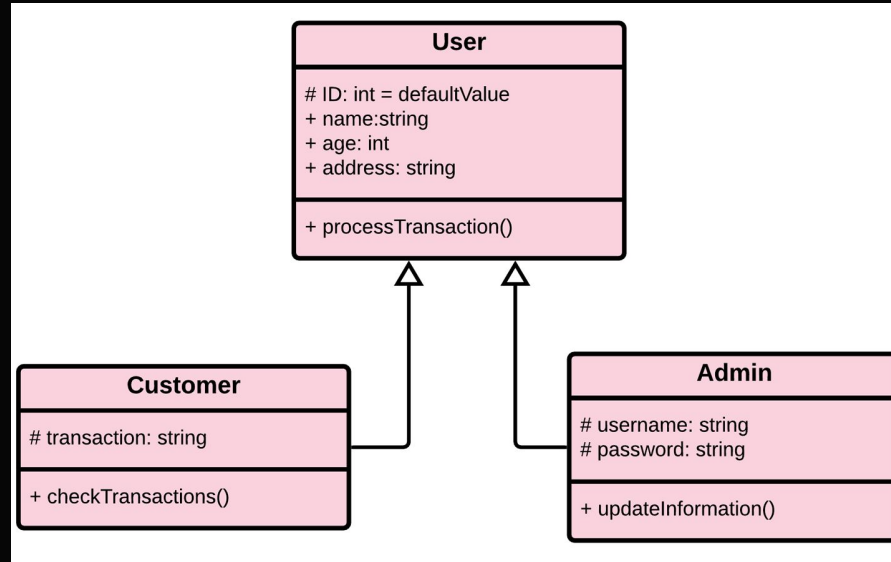
```
arr = [5, 1, 4, 2, 8]
n = len(arr)
for i in range(n):
    # Last i elements are already sorted
    for j in range(0, n - i - 1):
        # Swap if the current element is greater than the next
        if arr[j] > arr[j + 1]:
            arr[j], arr[j + 1] = arr[j + 1], arr[j]
print(arr) # Output: [1, 2, 4, 5, 8]

arr = [6, 3, 10, 1, 2]
n = len(arr)
for i in range(n):
    # Last i elements are already sorted
    for j in range(0, n - i - 1):
        # Swap if the current element is greater than the next
        if arr[j] > arr[j + 1]:
            arr[j], arr[j + 1] = arr[j + 1], arr[j]
print(arr) # Output: [1, 2, 3, 6, 10]

arr = [2, 3, 1, 8, 4]
n = len(arr)
for i in range(n):
    # Last i elements are already sorted
    for j in range(0, n - i - 1):
        # Swap if the current element is greater than the next
        if arr[j] > arr[j + 1]:
            arr[j], arr[j + 1] = arr[j + 1], arr[j]
print(arr) # Output: [1, 2, 3, 4, 8]
```

Object Oriented Programming

- A way to make custom data types that have functions and variables that they own



Object Oriented Programming (Example)

```
# This defines a class called Dog
class Dog:
    # The __init__ method runs when we create a new Dog
    def __init__(self, name, age):
        self.name = name # Each dog has a name
        self.age = age   # Each dog has an age

    # A method to make the dog bark
    def bark(self):
        print(f"{self.name} says woof!")

    # A method to show the dog's information
    def show_info(self):
        print(f"This is {self.name} and they are {self.age} years old.")


# Creating (instantiating) two dogs
dog1 = Dog("Buddy", 3)
dog2 = Dog("Luna", 5)

# Calling methods
dog1.bark()          # Output: Buddy says woof!
dog2.show_info()     # Output: This is Luna and they are 5 years old.
```

Python for (Deep) Reinforcement Learning

NumPy

- NumPy is the backbone of most major machine learning libraries and frameworks due to its incredibly fast implementations of matrix and statistics math
- Given machine learning's extreme amount of matrix operations it's incredibly important to be able to do these operations fast
 - e.g . training chatGPT can reach trillions to quadrillions of matrix calculations (1,000,000,000,000 - 1,000,000,000,000,000)

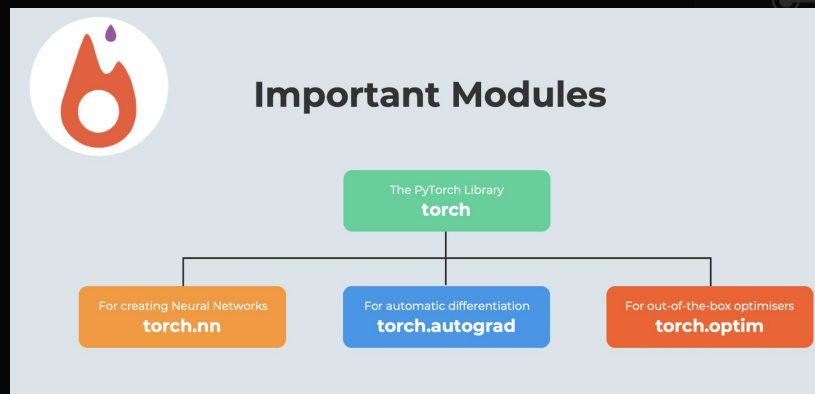


NumPy - Main Functionality

- Efficient storage and manipulation of large multidimensional arrays (ndarray)
- Supports fast element-wise operations
- Built-in fast operators for matrix math (dot product, mean, add, multiply, etc)
- Enables vectorized computations, reducing the need for explicit loops
- Offers advanced indexing, slicing, and reshaping for flexible data access
- Seamlessly integrates with other scientific libraries (e.g., SciPy, Pandas, TensorFlow)

PyTorch

- With PyTorch, you can build an entire neural network with a single line of code
- Holds efficient implementations of almost every machine learning algorithm you can think of, both classical and modern
- Builds on top of numpy to implement those models

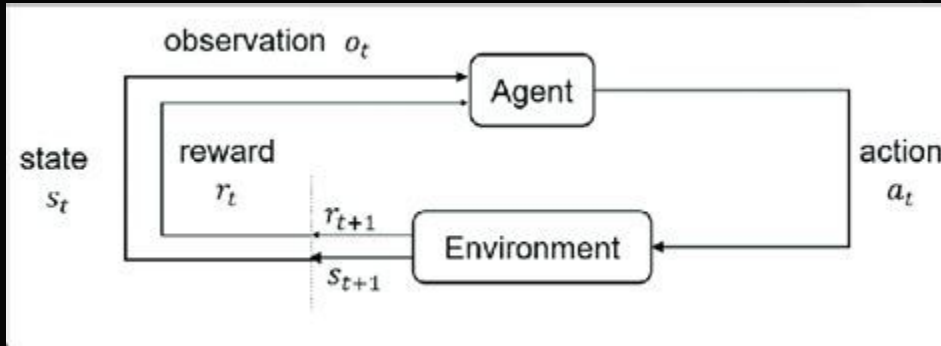


PyTorch - Main Functionalities

- Tensors for efficient multidimensional array operations with optional GPU acceleration
- Automatic differentiation for gradient computation
- A modular neural network API (`torch.nn`) for building and training deep learning models easily
- Built-in optimizers (e.g., SGD, Adam) for updating model parameters
- Easy GPU and CPU interoperability for data and model computation
- Strong ecosystem integration with libraries for vision, audio, and natural language processing

Gymnasium

- Gymnasium is an open source library that makes experimenting with Reinforcement Learning easy
- Has built in environments, models, and more to let you set up a full RL experiment in under 20 lines of code



Gymnasium Examples

Custom Environment

- Declare your action space
 - What kind of actions are you expecting?
 - E.g. one value from -100 to 100 representing left or right turn intensity
- Declare your observation space
 - What kind of observations should the agent be able to see to learn from?
 - E.g. how close the agent is to the landing zone in meters
- Defining a reward function
 - Define a reward function that lets the agent know when they're doing something good.
 - E.g. more reward for staying upright in a cartPole environment

Custom Agent

- What algorithm are you using to train your agent?
 - Many, many options:
 - DQN
 - DDPG
 - TD3
 - SAC
 - etc...

How to Debug on the Fly